



1. Feedforward neural networks

1.1. The perceptron


Manel Martínez-Ramón

- Inspired in the structure of the nervous system (McCulloch and Pitts, 1943).
- Described as an element with two possible states (0, 1).



- If a linear combination of the inputs (dendrites) is above a threshold, the output (axon) will be activated.

- The concept of artificial neural network was introduced by Rosemblatt (1958)



- First stage (retina): collects the observation.
- Second stage (projection): extracts the information.
- Third stage (association): maps the features into a set of responses.

- Roseblatt proved that his structure could learn from data.
- He developed the Mark 1 perceptron machine.
 - The Mark 1 was an electromechanical *learning machine*.
 - It had a camera of 400 pixels and attenuators driven by motors.
- The machine was able to classify *linearly separable* patterns.

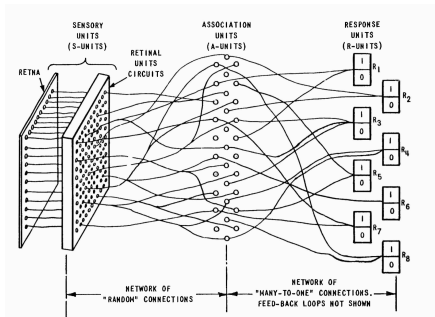



Figure 1 ORGANIZATION OF THE MARK I PERCEPTOR

Mark I Perceptron Operators' Manual, Cornell Aeronautical Laboratory (1960).

- A simplification of the perceptron is a binary classifier.
- Assume a set of observations in a column vector $\mathbf{x} = \{x_1 \cdots x_D\}^\top$, that can be arbitrarily labelled as “black” (-1) and “white” (+1) classes.




A classification function can be constructed from a *separating hyperplane* between both classes:

$$\mathbf{w}^\top \mathbf{x} + b = 0 \quad (1)$$

The classifier is defined as

$$f(\mathbf{x}) = \text{sign} \left(\mathbf{w}^\top \mathbf{x} + b \right) \quad (2)$$

as we will prove next.





- Vector \mathbf{w} is normal to the hyperplane defined by $\mathbf{w}^\top \mathbf{x} + b = 0$.
- Vector \mathbf{x}_0 is inside the plane, hence $\mathbf{w}^\top \mathbf{x}_0 + b = 0$
- Vector \mathbf{x}_1 can be written as $\mathbf{x}_1 = \mathbf{x}_0 + \mathbf{z}$
- Therefore, since $\angle \overline{\mathbf{x}_0 \mathbf{z}} > 90^\circ$:

$$\begin{aligned} f(\mathbf{x}_1) &= \text{sign}(\mathbf{w}^\top \mathbf{x}_1 + b) = \text{sign}(\mathbf{w}^\top (\mathbf{x}_0 + \mathbf{z}) + b) \\ &= \text{sign}(\mathbf{w}^\top \mathbf{x}_0 + b + \mathbf{w}^\top \mathbf{z}) = \text{sign}(\mathbf{w}^\top \mathbf{z}) = -1 \end{aligned} \tag{3}$$

Graphically, the classifier is a structure weights that represent the dot product and a bias, and a generic *activation function* σ :

$$f(\mathbf{x}) = \sigma \left(\mathbf{w}^\top \mathbf{x} + b \right) = \sigma \left(\sum_d w_d x_d + b \right)$$



The separating hyperplane in an example of 2 dimensions (left), and a representation of the classification function for the case of 3 dimensions.

- The perceptron rule is the first algorithm to train a learning machine, that was able to classify linearly separable patterns, and it was introduced by Rosenblatt in 1958.
- Let us start by defining the concept of error.

$$e_i = \frac{1}{2} \left(y_i - \text{sign}(\mathbf{w}^\top \mathbf{x}_i + b) \right) \quad (4)$$

This is simply 1 if the sample has been misclassified, and zero otherwise.

- The training criterion is simply “If a training sample is misclassified, then slightly move the boundary so it is properly classified”.

Assume a sample \mathbf{x}_k with label y_k and a classifier
 $f(\mathbf{x}) = \text{sign}(\mathbf{w}^\top \mathbf{x}_k + b) \neq y_k$ this is, the sample is misclassified.

- If the label is $y_k = 1$ and the classification is wrong, then

$$\mathbf{w}^{(k)\top} \mathbf{x}_k + b^{(k)} < 0$$

- The training rule is

$$\begin{aligned}\mathbf{w}^{(k+1)} &= \mathbf{w}^{(k)} + y_k \mathbf{x}_k \\ b^{(k+1)} &= b^{(k)} + y_k\end{aligned}\tag{5}$$

and, after the update

$$\left(\mathbf{w}^{(k)} + \mathbf{x}_k\right)^\top \mathbf{x}_k + b + 1 = \mathbf{w}^{(k)\top} \mathbf{x}_k + b^{(k)} + \|\mathbf{x}_k\|^2 + 1 > \mathbf{w}^{(k)\top} \mathbf{x}_k + b^{(k)}$$

- If the label is $y_i = -1$ and the classification is wrong, then

$$\mathbf{w}^\top \mathbf{x}_i + b > 0$$

and, after the update





$$(\mathbf{w} - \mathbf{x}_i)^\top \mathbf{x}_i + b - 1 < \mathbf{w}^\top \mathbf{x}_i + b$$

In both cases, the response of the classifier gets closer to the correct classification.

Theorem

If the training dataset is linearly separable, then the perceptron rule converges in a finite number of iterations (Novikoff, 1963).

Example of the perceptron rule



Example of the application of the perceptron rule in a set of separable data in dimension 2.

- The perceptron can be generalized to multiclass classification in a straightforward way to obtain the complete original algorithm.
- The algorithm presents two main limitations
 - The structure is purely linear.
 - The algorithm will not converge if the data is not linearly separable (i. e. if there are overlapping between classes).
- The second limitation can be overcome if the thresholded output (0, 1) is changed by a *soft* output that allows a more parsimonious update.
- Such an algorithm was not implemented in the Mark 1 because of the limitations of the electromechanical machine. For that purpose, an arithmetic unit is needed.

- In order to change the hard decision function $\text{sign}(\cdot)$ by a soft one, we can take the simplest choice, which is the linear one:

$$f(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + b$$

- The training criterion is then (MMSE)

$$\min_{\mathbf{w}, b} \mathbb{E} [e_i^2] = \min \mathbb{E} \left[\left(y_i - \mathbf{w}^\top \mathbf{x}_i - b \right)^2 \right] \quad (6)$$

Of course, the actual expectation cannot be computed, because the probability density functions of the random variables are not available, so the expectation will be approximated by a sample average.

- Assuming that N labelled samples \mathbf{x}_i, y_i are available, we introduce here matrix $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_N]$ and vector $\mathbf{y} = [y_1, \dots, y_N]$ containing all training samples and labels.
- By computing the gradients of the error with respect to \mathbf{w} and b and nulling them we obtain

$$\begin{aligned}\mathbf{w} &= \left(\mathbf{X}\mathbf{X}^\top\right)^{-1} \mathbf{X} (\mathbf{y} - b\mathbf{1}) \\ b &= \frac{1}{N} \sum_{i=1}^N \left(y_i - \mathbf{w}^\top \mathbf{x}_i\right)\end{aligned}\tag{7}$$

where $\mathbf{1}$ is a column of N ones. (Derivation as an exercise).

The above solution is too cumbersome, but if we extend both the input data and the weight vector, the solution is more compact. The extension is


$$\mathbf{x} \rightarrow \begin{bmatrix} \mathbf{x} \\ 1 \end{bmatrix}, \quad \mathbf{w} \rightarrow \begin{bmatrix} \mathbf{w} \\ b \end{bmatrix}$$

In this case, nulling the gradient gives

$$\mathbf{w} = \left(\mathbf{X}\mathbf{X}^\top \right)^{-1} \mathbf{X}\mathbf{y} \quad (8)$$

(derivation as an exercise) which is a compact solution, i.e. no iterations needed.

Here we derive a recursive solution and we compare it to the Perceptron rule. The method is based on a *gradient descent* approach: compute the gradient of the error wrt \mathbf{w} and move the weights in its opposite direction.



$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} - \mu \nabla_{\mathbf{w}} \mathbb{E}[e^2] \quad (9)$$

where

$$\nabla_{\mathbf{w}} \mathbb{E}[e^2] = \mathbf{X}\mathbf{X}^T \mathbf{w} - \mathbf{X}\mathbf{y} \quad (10)$$

Now we approximate Eq. (10) by using a single sample:

$$\begin{aligned}\nabla_{\mathbf{w}} \mathbb{E} [e^2] &= \mathbf{X}\mathbf{X}^\top \mathbf{w} - \mathbf{X}\mathbf{y} \\ &\approx \mathbf{x}_k \mathbf{x}_k^\top \mathbf{w} - \mathbf{x}_k y_k \\ &= \mathbf{x}_k \left(\mathbf{x}_k^\top \mathbf{w} - y_k \right) \\ &= -e_k \mathbf{x}_k\end{aligned}\tag{11}$$

Where $e_k = y_k - \mathbf{x}_k^\top \mathbf{w}$. This leads to the following update rule

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} + \mu e_k \mathbf{x}_k\tag{12}$$



Compare the LMS solution with the Perceptron rule in Eq. (5). The error as defined for LMS is $e_k = y_k - \mathbf{x}_k^\top \mathbf{w} = y_k - \mathbf{w}^\top \mathbf{x}_k$, but in the perceptron rule, the error is defined as

$$e_k = \frac{1}{2} \left(y_k - \text{sign}(\mathbf{w}^\top \mathbf{x}_k + b) \right) = \begin{cases} y_k, & \mathbf{x}_k \text{ misclassified} \\ 0, & \text{otherwise.} \end{cases}$$

If we use this error and $\mu = 1$ the LMS then becomes the perceptron rule

$$\mathbf{w}^{(k+1)} = \begin{cases} \mathbf{w}^{(k)} + y_i \mathbf{x}_i, & \mathbf{x}_k \text{ misclassified} \\ 0, & \text{otherwise.} \end{cases}$$

- The perceptron uses a hard output (sign detection) as an output, that translates the *affine* transformation $\mathbf{w}^T \mathbf{x} + b$ into one of two states $(-1, +1)$.
- In order to implement the LMS, we just remove the sign detection.
- A *sigmoid* function can be used to produce an output that can be interpreted as a soft state or a probability of a state.



Hyperbolic tangent function (left) and logistic function.

- The logistic function or the hyperbolic tangent are classic neural network activations, but nowadays, the logistic is used in combination with other functions that will be studied further.
- The hyperbolic tangent, the logistic function and their derivatives are:

$$\begin{aligned}\tanh(z) &= \frac{e^z - e^{-z}}{e^z + e^{-z}}, & \frac{d}{dz} \tanh(z) &= 1 - \tanh^2(z) \\ \sigma(z) &= \frac{1}{1 + e^{-z}}, & \frac{d}{dz} \sigma(z) &= \sigma(z) (1 - \sigma(z))\end{aligned}\tag{13}$$

Now, if the output of the classifier is

$$f(\mathbf{x}) = \tanh \left(\mathbf{w}^\top \mathbf{x}_i + b \right) \quad (14)$$





then the criterion to optimize the parameters will be

$$\min_{\mathbf{w}, b} \mathbb{E} [e_i^2] \approx \min_{\mathbf{w}, b} \sum_{i=1}^N \left(y_i - \tanh \left(\mathbf{w}^\top \mathbf{x}_i + b \right) \right)^2 \quad (15)$$

which leads to the update rule

$$\begin{aligned} \mathbf{w}^{(k+1)} &= \mathbf{w}^{(k)} + \mu \sum_{i=1}^N e_i (1 - f^2(\mathbf{x}_i)) \mathbf{x}_i \\ b^{(k+1)} &= b^{(k)} + \mu \sum_{i=1}^N e_i (1 - f^2(\mathbf{x}_i)) \end{aligned} \quad (16)$$

The proof is left as an exercise.



Example of the application of the MMSE criterion to a perceptron with hyperbolic tangent activation. The first row corresponds to a separable case and the second one to a non separable one. In both cases the algorithm converges to a solution.